

Concurrency

Reading (two lectures)

Chapter 14, except section 14.3 and except pages 461-464

JSR 133 (Java Memory Model) FAQ

Note: book presentation of memory model is obsolete

Outline

- ➔ • What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example

Concurrency

Two or more sequences of events occur in parallel

Note:

Process: sequential program running on a processor

The promise of concurrency

- Speed
 - If a task takes time t on one processor, shouldn't it take time t/n on n processors?
- Availability
 - If one process is busy, another may be ready to help
- Distribution
 - Processors in different locations can collaborate to solve a problem or work together
- Humans do it so why can't computers?
 - Vision, cognition appear to be highly parallel activities

Concurrency on machines

- Multiprogramming

- A single computer runs several programs at the same time
- Each program proceeds sequentially
- Actions of one program may occur between two steps of another

- Multiprocessors

- Two or more processors may be connected
- Programs on one processor communicate with programs on another
- Actions may happen simultaneously

Challenges

- Concurrent programs are harder to get right
 - Folklore: Need at least an order of magnitude in speedup for concurrent prog to be worth the effort
- Some problems are inherently sequential
 - Theory – circuit evaluation is P-complete
 - Practice – many problems need coordination and communication among sub-problems
- Specific issues
 - Communication – send or receive information
 - Synchronization – wait for another process to act
 - Atomicity – do not stop in the middle and leave a mess

Basic question for this course

- How can programming languages make concurrent programming easier?
- Which abstractions are most effective?
 - What are the advantages and disadvantages of various approaches?

Apart from basic concepts, this lecture covers past ideas (cobegin / coend, actor model) and current Java.

Next week we look at forward-looking research ideas.

Outline

- What is concurrency?
- ➔ • Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example


Basic issue: race conditions

- Sample action

```
procedure sign_up(person)
begin
  index := index + 1;
  list[index] := person;
end;
```

- Problem with parallel execution

```
sign_up(fred) || sign_up(bill);
```



Resolving conflict between processes

- Critical section
 - Two processes may access shared resource
 - Inconsistent behavior if two actions are interleaved
 - Allow only one process in *critical section*
- Potential solution: Locks?
- Problem: Deadlock
 - Process may hold some locks while awaiting others
 - *Deadlock* occurs when no process can proceed

Locks and Waiting

<initialize concurrency control>

Thread 1:

<wait>

sign_up(fred); // critical section

<signal>

Thread 2:

<wait>

sign_up(bill); // critical section

<signal>

Need atomic operations to implement wait


Mutual exclusion primitives

- Atomic test-and-set
 - Instruction atomically reads and writes some location
 - Common hardware instruction
 - Used to implement a busy-waiting loop to get mutual exclusion
- Semaphore
 - Avoid busy-waiting loop
 - Keep queue of waiting processes
 - Scheduler has access to semaphore; process sleeps
 - Disable interrupts during semaphore operations
 - OK since operations are short

State of the art

- Concurrent programming is difficult
 - Race conditions, deadlock are pervasive
- Languages should be able to help
 - Capture useful paradigms, patterns, abstractions
- Other tools are needed
 - Testing is difficult for multi-threaded programs
 - Many race-condition detectors being built today
 - Static detection: conservative, may be too restrictive
 - Run-time detection: may be more practical for now

Outline

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
-  • Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example

Cobegin/coend

- Limited concurrency primitive
- Example

`x := 0;`

`cobegin`


`begin x := 1; x := x+1 end;`

`begin x := 2; x := x+1 end;`

`coend;`

`print(x);`

} execute sequential
blocks in parallel



Atomicity at level of assignment statement

Properties of cobegin/coend


- **Advantages**
 - Create concurrent processes
 - Communication: Shared variables
- **Limitations**
 - Mutual exclusion: none
 - Atomicity: none
 - Number of processes is fixed by program structure
 - Cannot abort processes
 - All must complete before parent process can go on

Actors

[Hewitt, Agha, Tokoro, Yonezawa, ...]

- Each actor (object) has a script
- In response to input, actor may atomically
 - create new actors
 - initiate communication
 - change internal state
- Communication is
 - Buffered, so no message is lost
 - Guaranteed to arrive, but not in sending order
 - Order-preserving communication is harder to implement
 - Programmer can build ordered primitive from unordered
 - Inefficient to have ordered communication when not needed

Example



Actor program

- Stack node $\overbrace{\text{parameters}}$
a stack_node with acquaintances content and link
if operation requested is a pop and content != nil then
become forwarder to link
send content to customer
if operation requested is push(new_content) then
let P=new stack_node with current acquaintances (a clone)
become stack_node with acquaintances new_content and P

Hard to read but it does the “obvious” thing, except that the concept of *forwarder* is unusual....

Forwarder

- Stack before pop




- Stack after pop



- Node “disappears” by becoming a forwarder node.
The system manages forwarded nodes in a way that makes them invisible to the program.
(Exact mechanism doesn’t matter)

Concurrency

- Several actors may operate concurrently



- Concurrency not controlled explicitly by program
 - Messages sent by one actor can be received and processed by others sequentially or concurrently

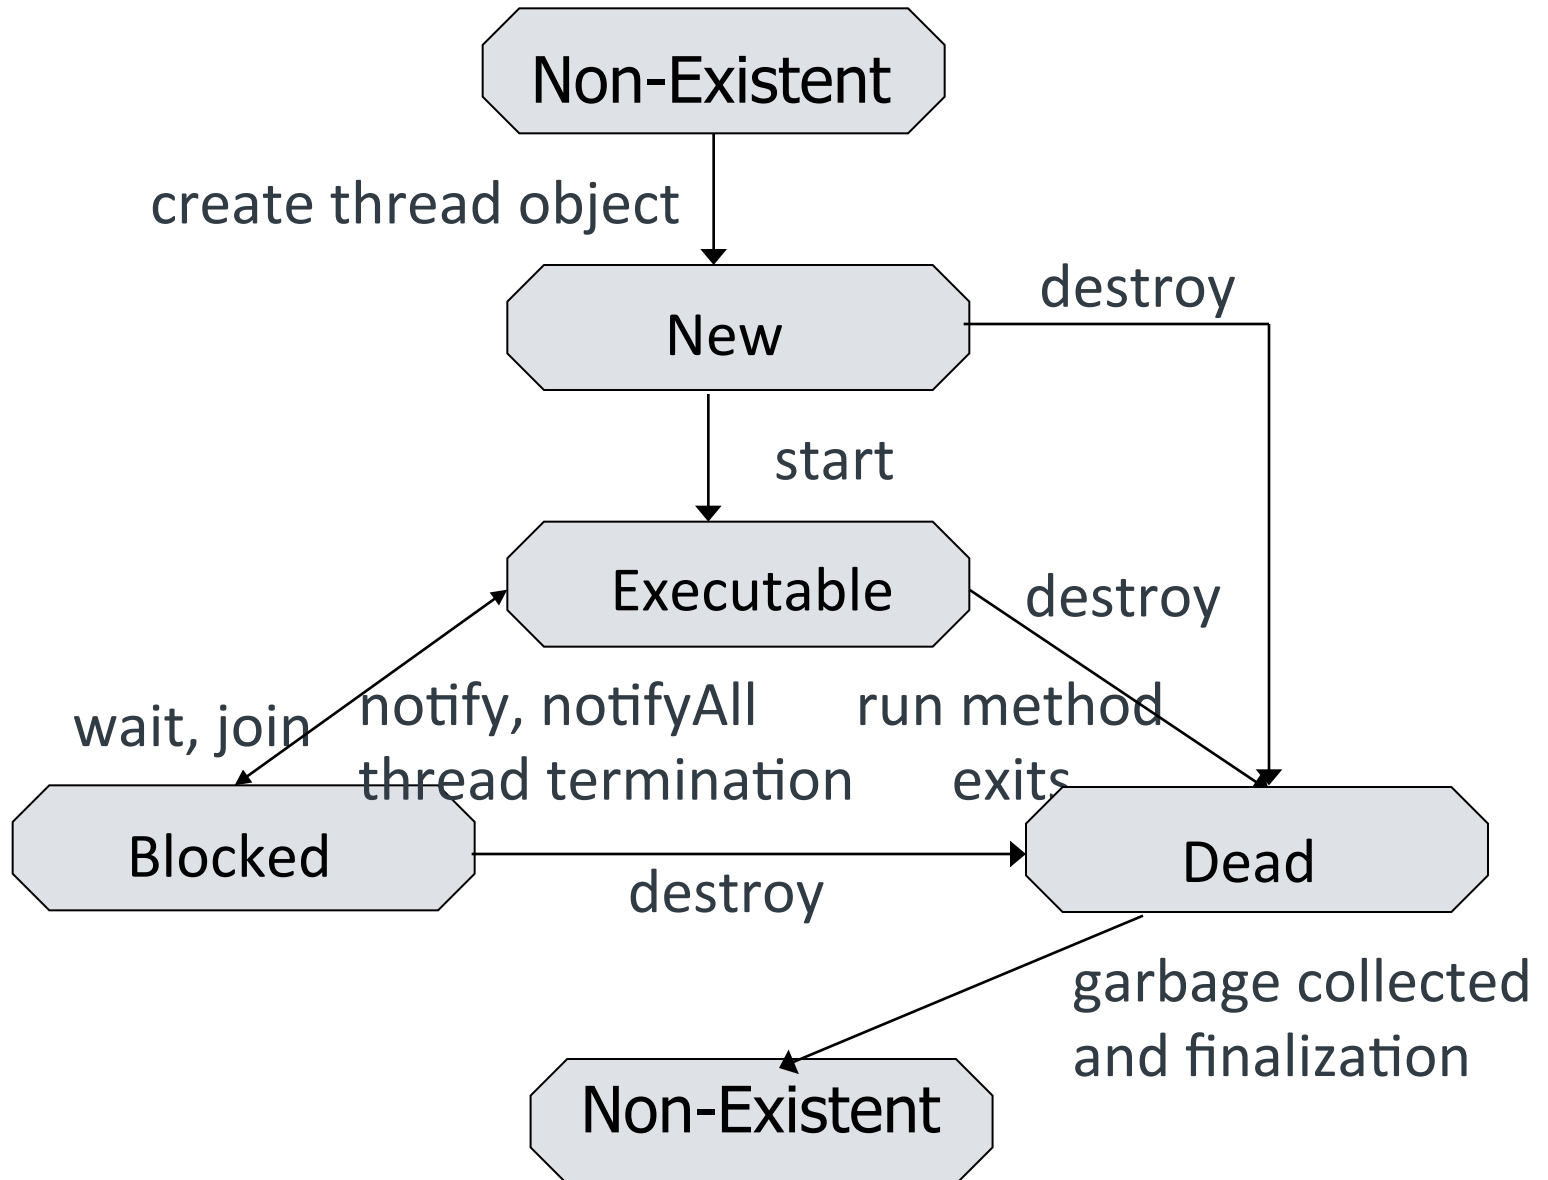
Outline

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- ➡ • Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example

Java Concurrency

- **Threads**
 - Create process by creating thread object
- **Communication**
 - Shared variables
 - Method calls
- **Mutual exclusion and synchronization**
 - Every object has a lock (inherited from class Object)
 - synchronized methods and blocks
 - Synchronization operations (inherited from class Object)
 - wait: pause current thread until another thread calls notify
 - notify: wake up waiting threads

Java Thread States




Interaction between threads

- Shared variables
 - Two threads may assign/read the same variable
 - Programmer responsibility
 - Avoid race conditions by explicit synchronization !!
- Method calls
 - Two threads may call methods on the same object
- Synchronization primitives
 - Each object has internal lock, inherited from Object
 - Synchronization primitives based on object locking

Synchronization

- Provides mutual exclusion
 - Two threads may have access to some object
 - If one calls a synchronized method, this locks object
 - If the other calls a synchronized method on same object, this thread blocks until object is unlocked



Synchronized methods

- Marked by keyword

`public synchronized void commitTransaction(...) {...}`


- Not part of method signature
 - sync method equivalent to unsync method with body consisting of a *synchronized block*
 - subclass may replace a synchronized method with unsynchronized method

Example

[Lea]

```
class LinkedCell {           // Lisp-style cons cell containing
    protected double value; // value and link to next cell
    protected final LinkedCell next;
    public LinkedCell (double v, LinkedCell t) {
        value = v; next = t;
    }
    public synchronized double getValue() {
        return value;
    }
    public synchronized void setValue(double v) {
        value = v; // assignment not atomic
    }
    public LinkedCell next() { // no synch needed
        return next;
    }
}
```

Producer-Consumer



- How do we do this in Java?

Solution to producer-consumer

- Basic idea
 - Consumer must **wait** until something is in the buffer
 - Producer must **notify** waiting consumers when item available
- More details
 - Consumer waits
 - While waiting, must *sleep* – accomplished with the wait method
 - Need condition recheck loop
 - Producer notifies
 - Must *wake up* at least one consumer
 - This is accomplished with the notify method

Stack<T>: produce, consume methods

```
public synchronized void produce (T object) {  
    stack.add(object);  
    notify();  
}
```

```
public synchronized T consume () {  
    while (stack.isEmpty()) {  
        try {  
            wait();  
        } catch (InterruptedException e) {  
        }  
    }  
    int lastElement = stack.size() - 1;  
    T object = stack.get(lastElement);  
    stack.remove(lastElement);  
    return object;  
}
```

Limitations of Java 1.4 primitives

- No way to back off from an attempt to acquire a lock
 - Cannot give up after waiting for a specified period of time
 - Cannot cancel a lock attempt after an interrupt
- No way to alter the semantics of a lock
 - Reentrancy, read versus write protection, fairness, ...
- No access control for synchronization
 - Any method can perform synchronized(obj) for any object
- Synchronization is done within methods and blocks
 - Limited to block-structured locking
 - Cannot acquire a lock in one method and release it in another

See <http://java.sun.com/developer/technicalArticles/J2SE/concurrency/>

Concurrency references

- Thread-safe classes
 - B Venners, Designing for Thread Safety, JavaWorld, July 1998: <http://www.artima.com/designtechniques/threadsafety.html>
- Nested monitor lockout problem
 - <http://www-128.ibm.com/developerworks/java/library/j-king.html?dwzone=java>
- Inheritance anomaly
 - G Milicia, V Sassone: The Inheritance Anomaly: Ten Years After, SAC 2004: <http://citeseer.ist.psu.edu/647054.html>
- Java memory model
 - See <http://www.cs.umd.edu/~jmanson/java.html>
 - and <http://www.cs.umd.edu/users/jmanson/java/journal.pdf>
- Race conditions and correctness
 - See slides: lockset, vector-clock algorithms
- Atomicity and tools
 - See <http://www.cs.uoregon.edu/activities/summerschool/summer06/>

More detail in references than required by course

Outline

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - ➡ – Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example

Thread safety

- Concept
 - The fields of an object or class always maintain a valid state, as observed by other objects and classes, even when used concurrently by multiple threads
- Why is this important?
 - Classes designed so each method preserves state invariants
 - Example: priority queues represented as sorted lists
 - Invariants hold on method entry and exit
 - If invariants fail in the middle of execution of a method, then concurrent execution of another method call will observe an inconsistent state (state where the invariant fails)
 - What's a “valid state”? Serializability ...

Example

(two slides)

```
public class RGBColor {
    private int r;  private int g;  private int b;
    public RGBColor(int r, int g, int b) {
        checkRGBVals(r, g, b);
        this.r = r;  this.g = g;  this.b = b;
    }

    ...

    private static void checkRGBVals(int r, int g, int b) {
        if (r < 0 || r > 255 || g < 0 || g > 255 ||
            b < 0 || b > 255) {
            throw new IllegalArgumentException();
        }
    }
}
```

Example

(continued)

```
public void setColor(int r, int g, int b) {  
    checkRGBVals(r, g, b);  
    this.r = r;    this.g = g;    this.b = b;  
}
```

```
public int[] getColor() { // returns array of three ints: R, G, and B  
    int[] retVal = new int[3];  
    retVal[0] = r;    retVal[1] = g;    retVal[2] = b;  
    return retVal;  
}
```

```
public void invert() {  
    r = 255 - r;    g = 255 - g;    b = 255 - b;  
}
```

Question: what goes wrong with multi-threaded use of this class?

Some issues with RGB class

- Read/write conflicts
 - If one thread reads while another writes, the color that is read may not match the color before *or* after
- Write/write conflicts
 - If two threads try to write different colors, result may be a “mix” of R,G,B from two different colors

How to make classes thread-safe

1. Synchronize critical sections

- Make fields private
- Synchronize sections that should not run concurrently

2. Make objects immutable

- State cannot be changed after object is created

```
public RGBColor invert() {  
    RGBColor retVal = new RGBColor(255 - r, 255 - g, 255 - b);  
    return retVal;  
}
```

- Use pure functional programming for concurrency

3. Use a thread-safe wrapper

- See next slide ...

How to make classes thread-safe: thread-safe wrapper

- Idea
 - New thread-safe class has objects of original class as fields
 - Wrapper class provides methods to access original class object

- Example

```
public synchronized void setColor(int r, int g, int b) {  
    color.setColor(r, g, b);  
}  
public synchronized int[] getColor() {  
    return color.getColor();  
}  
public synchronized void invert() {  
    color.invert();  
}
```


Comparison

- Synchronizing critical sections
 - Good default approach for building thread-safe classes
 - Only way to allow wait() and notify()
- Using immutable objects
 - Good if objects are small, simple abstract data type
 - Benefit: pass to methods without alias issues, unexpected side effects
 - Examples: Java String and primitive type wrappers Integer, Long, Float, etc.
- Using wrapper objects
 - Can give clients choice between thread-safe version and non-safe
 - Works with existing class that is not thread-safe
 - Example: Java 1.2 collections library – classes are not thread safe but some have class method to enclose objects in thread-safe wrapper

Performance issues

- Why not just synchronize everything?
 - Performance costs
 - Possible risks of deadlock from too much locking
- Performance in current Sun JVM
 - Synchronized methods are 4 to 6 times slower than non-synchronized
- Performance in general
 - Unnecessary blocking and unblocking of threads can reduce concurrency
 - Immutable objects can be short-lived, increase garbage collector

Outline

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - ➡ – Java memory model
 - Concurrent hash map example

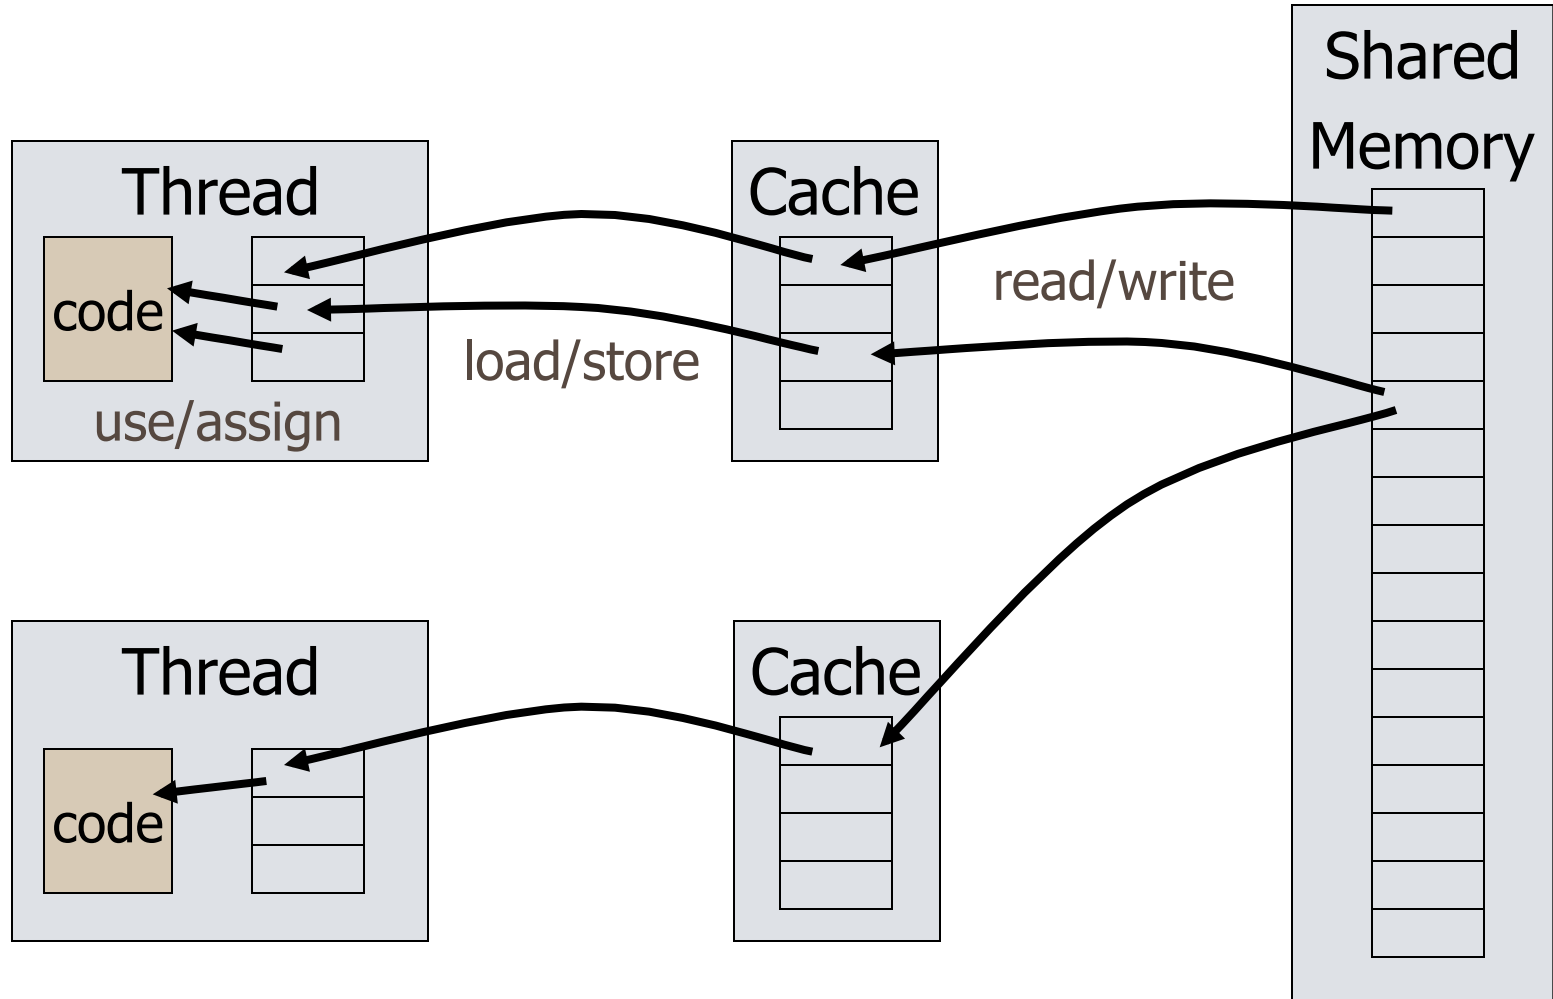
Java Memory Model

- Semantics of multithreaded access to shared memory
 - Competitive threads access shared data
 - Can lead to data corruption
 - Need semantics for incorrectly synchronized programs
- Determines
 - Which program transformations are allowed
 - Should not be too restrictive
 - Which program outputs may occur on correct implementation
 - Should not be too generous

Reference:

<http://www.cs.umd.edu/users/pugh/java/memoryModel/jsr-133-faq.html>

Memory Hierarchy




Old memory model placed complex constraints on read, load, store, etc.

Program and locking order

Thread 1

Thread 2



Race conditions

- “Happens-before” order
 - Transitive closure of program order and synchronizes-with order
- Conflict
 - An *access* is a read or a write
 - Two accesses *conflict* if at least one is a write
- Race condition
 - Two accesses form a *data race* if they are from different threads, they conflict, and they are not ordered by happens-before

Race cond

Subtle issue: program order
as written, or as compiled
and optimized ?

- “Happens-before” order
 - Transitive closure of program order and synchronizes-with order
- Conflict
 - An *access* is a read or a write
 - Two accesses *conflict* if at least one is a write
- Race condition
 - Two accesses form a *data race* if they are from different threads, they conflict, and they are not ordered by happens-before


Memory Model Question

- How should the compiler and run-time system be allowed to schedule instructions?
- Possible partial answer
 - If instruction A occurs in Thread 1 before release of lock, and B occurs in Thread 2 after acquire of same lock, then A must be scheduled before B
- Does this solve the problem?
 - Too restrictive: if we prevent reordering in Thread 1,2
 - Too permissive: if arbitrary reordering in threads
 - Compromise: allow local thread reordering that would be OK for sequential programs

Instruction order and serializability

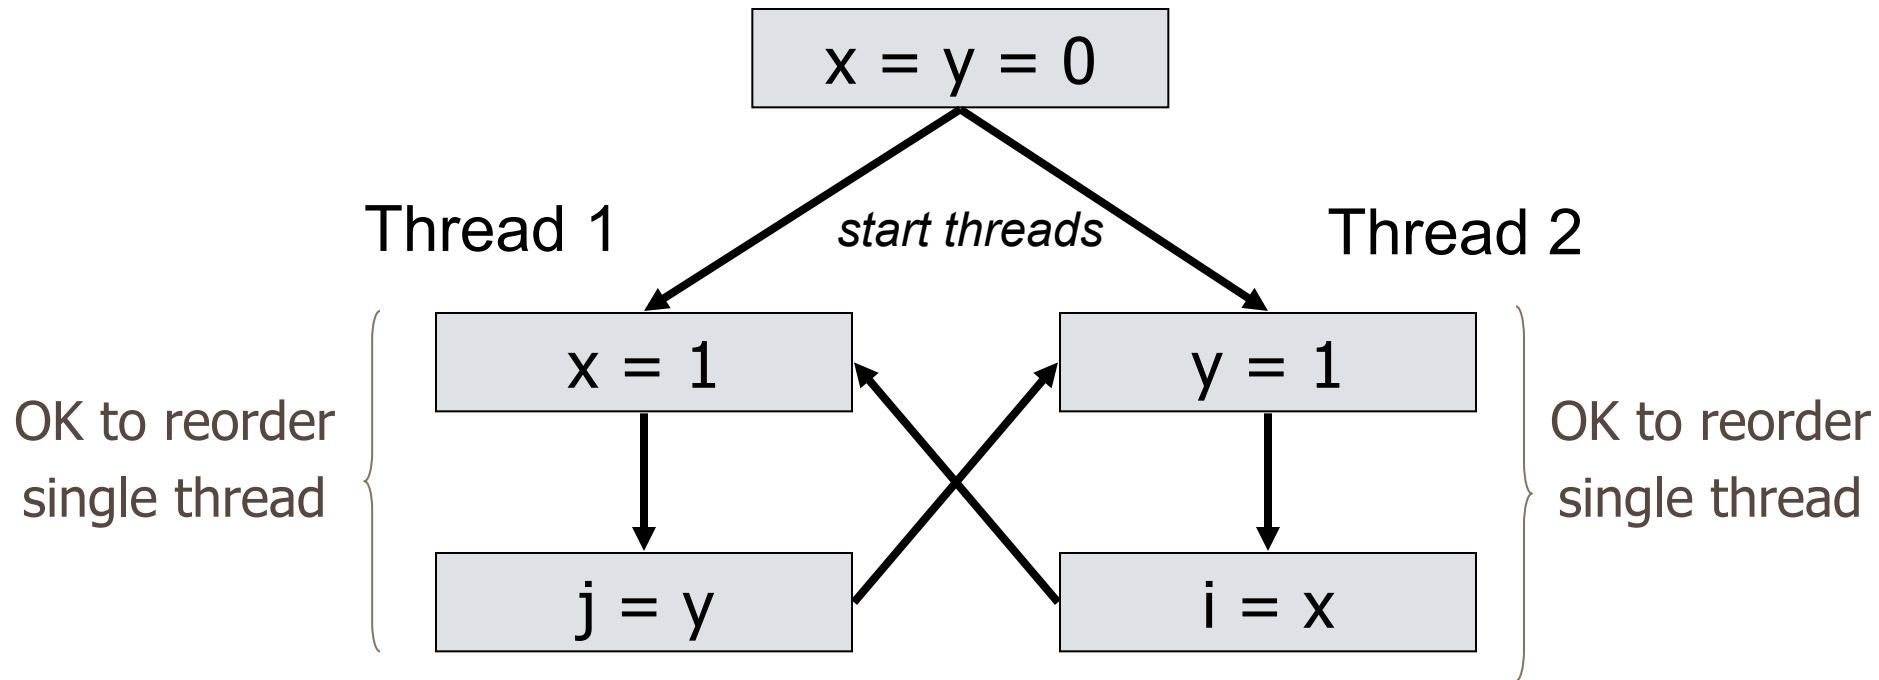
- Compilers can reorder instructions
 - If two instructions are independent, do in any order
 - Take advantage of registers, etc.
- Correctness for sequential programs
 - Observable behavior should be same as if program instructions were executed in the order written
- Sequential consistency for concurrent programs
 - If program P has no data races, then memory model should guarantee sequential consistency
 - Question: what about programs *with* races?
 - Much of complexity of memory model is for reasonable behavior for programs with races (need to test, debug, ...)

Example program with data race



Can we end up with $i = 0$ and $j = 0$?

Sequential reordering + data race



How can $i = 0$ and $j = 0$?


Java definition considers this OK since there is a data race

[Manson, Pugh]

Allowed sequential reordering

- “Roach motel” ordering
 - Compiler/processor can move accesses into synchronized blocks
 - Can only move them out under special circumstances, generally not observable
- Release only matters to a matching acquire
- Some special cases:
 - locks on thread local objects are a no-op
 - reentrant locks are a no-op
 - Java SE 6 (Mustang) does optimizations based on this

Something to prevent ...



- Must not result in $r1 = r2 = 42$
 - Imagine if 42 were a reference to an object!
- Value appears “out of thin air”
 - This is causality run amok
 - Legal under a simple “happens-before” model of possible behaviors

Summary of memory model

- Strong guarantees for race-free programs
 - Equivalent to interleaved execution that respects synchronization actions
 - Thread reordering must preserve sequential semantics of thread
- Weaker guarantees for programs with races
 - Allows program transformation and optimization
 - No weird out-of-the-blue program results
- Form of actual memory model definition
 - Happens-before memory model (examples on next slide)
 - Additional condition: for every action that occurs, there must be identifiable cause in the program

Happens-Before orderings

- Starting a thread happens-before the run method of the thread
- The termination of a thread happens-before a join with the terminated thread
- Volatile fields
- Many `util.concurrent` methods set up happen-before orderings
 - placing an object into any concurrent collection happens-before the access or removal of that element from the collection

Outline

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - ➡ – Concurrent hash map example

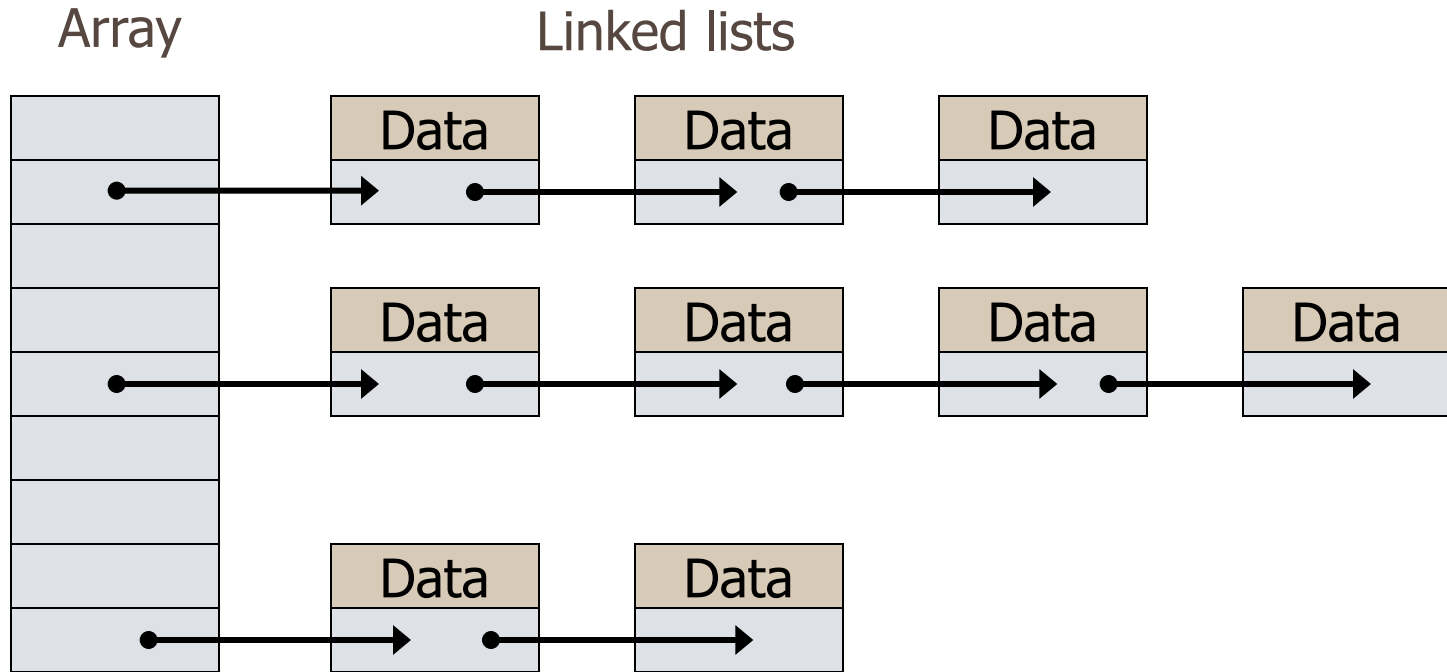
Example: Concurrent Hash Map

- Implements a hash table
 - Insert and retrieve data elements by key
 - Two items in same bucket placed in linked list
 - Allow read/write with minimal locking
- Tricky

“ConcurrentHashMap is both a very useful class for many concurrent applications and a fine example of a class that understands and exploits the subtle details of the Java Memory Model (JMM) to achieve higher performance. ... Use it, learn from it, enjoy it – but unless you're an expert on Java concurrency, you probably shouldn't try this on your own.”


See <http://www-106.ibm.com/developerworks/java/library/j-jtp08223>

ConcurrentHashMap



- Concurrent operations
 - read: no problem
 - read/write: OK if different lists
 - read/write to same list: clever tricks sometimes avoid locking

ConcurrentHashMap Tricks



- **Immutability**
 - List cells are immutable, except for data field
 - ⇒ read thread sees linked list, even if write in progress
- **Add to list**
 - Can cons to head of list, like Lisp lists
- **Remove from list**
 - Set data field to null, rebuild list to skip this cell
 - Unreachable cells eventually garbage collected

Summary

- What is concurrency?
- Basic issues in concurrency
 - Race conditions, locking, deadlock, mutual exclusion
- Simple language approaches (Past ideas)
 - Cobegin/Coend (Concurrent Pascal), Actor model
- Java Concurrency
 - Threads, synchronization, wait/notify
 - Methods for achieving thread safety
 - Java memory model
 - Concurrent hash map example

Happy Thanksgiving!

Two examples of problems in concurrent Java

Problem with language specification

- Java Lang Spec allows access to partial objects

```
class Broken {  
    private long x;  
    Broken() {  
        new Thread() {  
            public void run() { x = -1; }  
        }.start();  
        x = 0;  
    }  
}
```

Thread created within constructor can access the object not fully constructed

Nested monitor lockout problem

- Background: wait and locking
 - *wait* and *notify* used within synchronized code
 - Purpose: make sure that no other thread has called method of same object
 - *wait* within synchronized code causes the thread to give up its lock and sleep until notified
 - Allow another thread to obtain lock and continue processing
- Problem
 - Calling a blocking method within a synchronized method can lead to deadlock

Nested Monitor Lockout Example

```
class Stack {  
    LinkedList list = new LinkedList();  
    public synchronized void push(Object x) {  
        synchronized(list) {  
            list.addLast( x ); notify();  
        }  
    }  
    public synchronized Object pop() {  
        synchronized(list) {  
            if ( list.size() <= 0 ) wait();  
            return list.removeLast();  
        }  
    }  
}
```

Releases lock on Stack object but not lock on list;
a push from another thread will deadlock

Preventing nested monitor deadlock

- Two programming suggestions
 - No blocking calls in synchronized methods, or
 - Provide some nonsynchronized method of the blocking object
- No simple solution that works for all programming situations